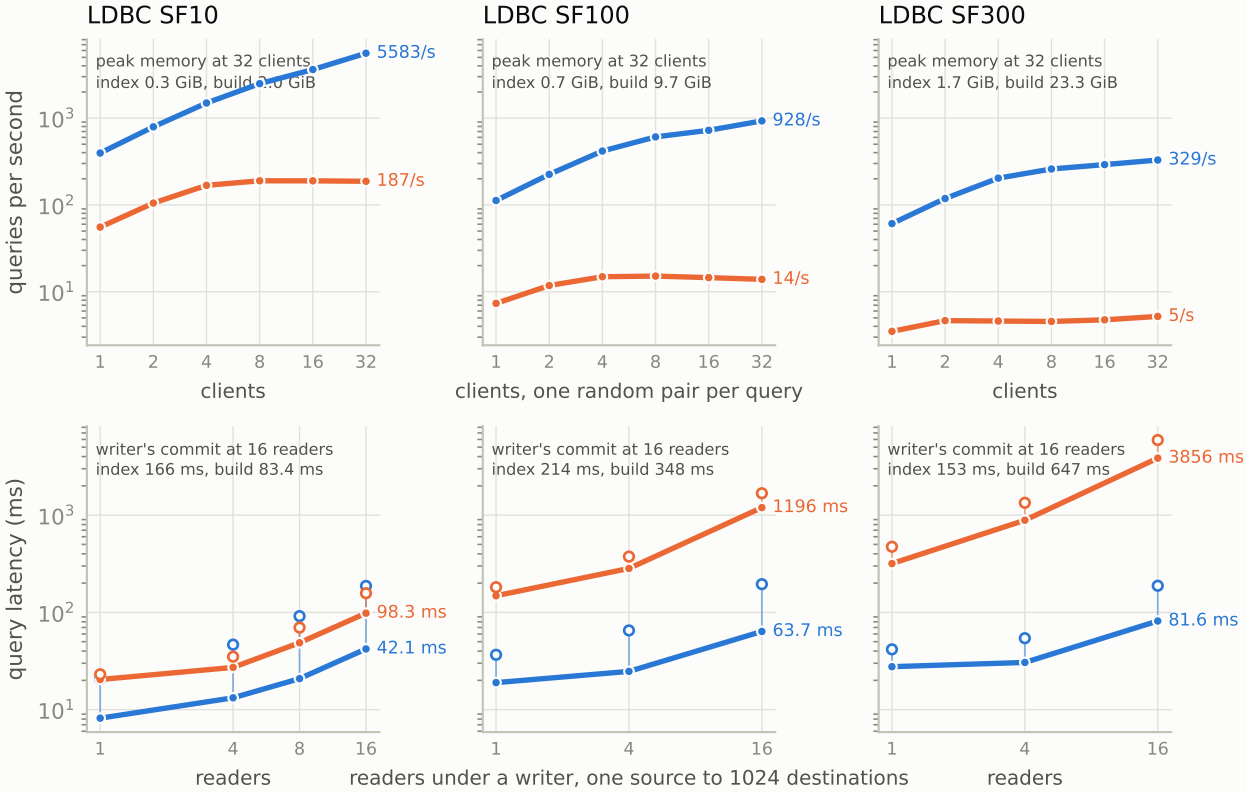
LDBC SF10
LDBC SF100
LDBC SF300
queries per second
peak memory at 32 clients
index 0.3 GiB, build 1.0 GiB
5583/s
187/s
peak memory at 32 clients
index 0.7 GiB, build 9.7 GiB
928/s
14/s
peak memory at 32 clients
index 1.7 GiB, build 23.3 GiB
329/s
5/s
clients
clients, one random pair per query
clients
query latency (ms)
writer's commit at 16 readers
index 166 ms, build 83.4 ms
98.3 ms
42.1 ms
writer's commit at 16 readers
index 214 ms, build 348 ms
1196 ms
63.7 ms
writer's commit at 16 readers
index 153 ms, build 647 ms
3856 ms
81.6 ms
readers
readers under a writer, one source to 1024 destinations
readers
DuckPGQ index
DuckPGQ build
p99 of the same queries