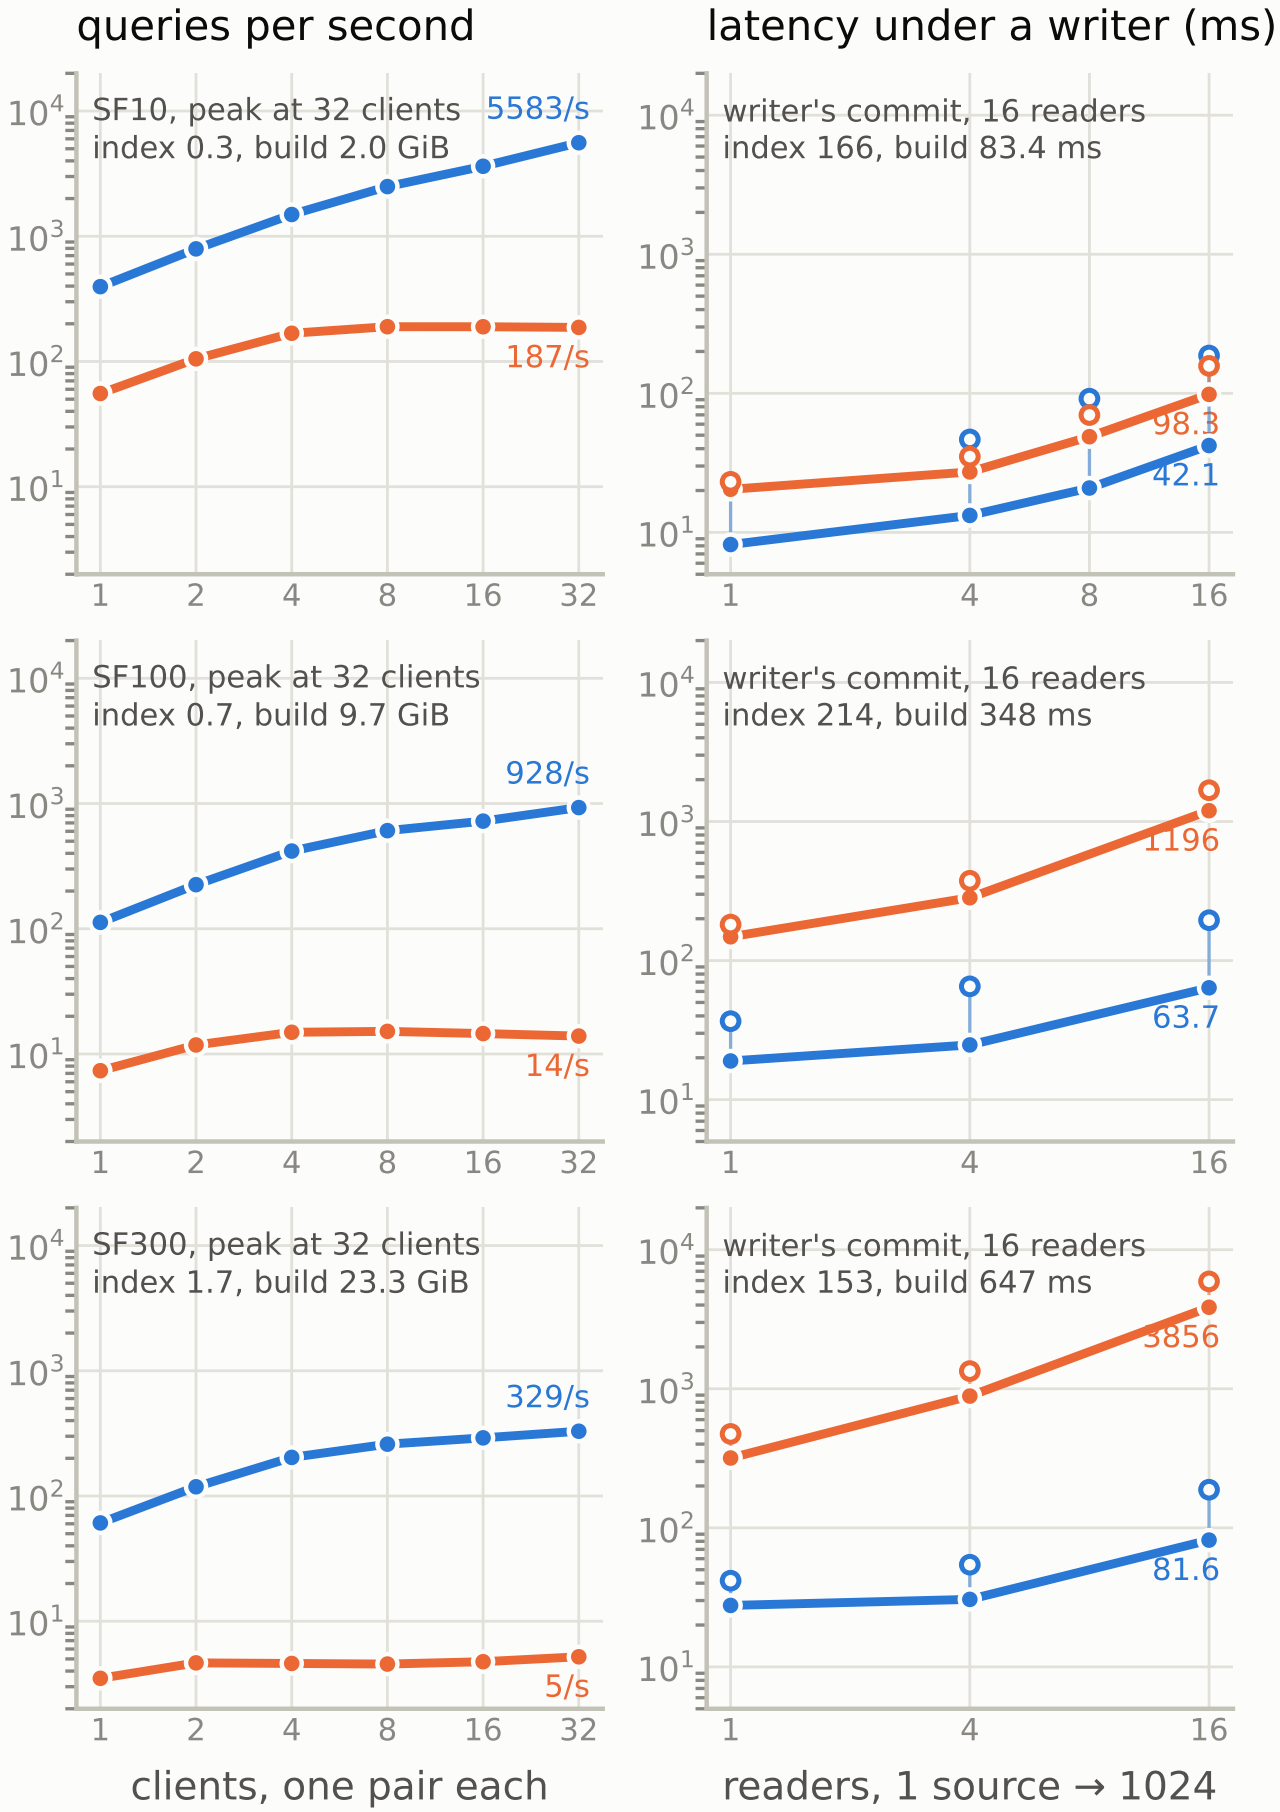
queries per second
latency under a writer (ms)
SF10, peak at 32 clients
index 0.3, build 2.0 GiB
5583/s
187/s
writer's commit, 16 readers
index 166, build 83.4 ms
98.3
42.1
SF100, peak at 32 clients
index 0.7, build 9.7 GiB
928/s
14/s
writer's commit, 16 readers
index 214, build 348 ms
1196
63.7
SF300, peak at 32 clients
index 1.7, build 23.3 GiB
329/s
5/s
writer's commit, 16 readers
index 153, build 647 ms
3856
81.6
clients, one pair each
readers, 1 source → 1024
DuckPGQ index
DuckPGQ build
p99